

# ☁️ PM2.5 Air Quality Prediction Dashboard

Real-time PM2.5 monitoring and 24-hour prediction system using Transformer neural network, deployed on Railway with auto-updating hourly crawler.

## 📋 Features

- **Real-time PM2.5 Monitoring** - Live data from Taiwan EPA API
- **24-Hour Predictions** - Transformer model forecasting
- **Interactive Visualizations** - Chart.js powered charts
- **AI Chatbot** - Gemini-powered RAG for historical queries
- **Auto-Update** - Hourly crawler + model inference
- **Modern UI** - Glassmorphism design, light mode

## 🏗 Architecture

- **Frontend:** HTML/CSS/JavaScript + Chart.js
- **Backend:** Flask + APScheduler
- **Database:** SQLite with Railway Volume (ALL data 2018-2025, ~61,000+ hours)
- **Model:** Transformer (uses last 720h → predicts 24h)
- **RAG:** Gemini API (queries full 2018-2025 history)
- **Deployment:** Railway with Docker

## 📁 Project Structure

```
pm25_dashboard/
├── backend/
│   ├── app.py      # Flask main app
│   ├── config.py   # Configuration
│   ├── database.py # SQLite utilities
│   ├── init_db.py  # Initialize DB from CSV
│   ├── crawler.py  # EPA API crawler
│   ├── prediction_service.py # Model inference
│   ├── rag_service.py # Gemini chatbot
│   └── scheduler.py  # APScheduler
└── frontend/
```

```
|   └── index.html
|   └── static/
|       ├── css/style.css
|       └── js/
|           ├── dashboard.js
|           └── chat.js
|   └── models/
|       └── best_model.keras # YOUR TRAINED MODEL
|   └── data/
|       └── pm25_data.db    # Historical database
└── Dockerfile
└── railway.toml
└── requirements.txt
└── README.md
```

## Deployment Instructions

### Step 1: Prepare Your Data & Model

1. Initialize database with historical data (run locally):

```
bash
python -m backend.init_db --csv /path/to/all_pm25_data.csv
```

Expected output:

```
 Data Range:
Start: 2018-01-01 00:00
End: 2025-11-23 14:00
Total: 61,320 hours (2,555 days)
```

```
 Database initialized successfully!
Total measurements: 61,320
Database size: ~8 MB
```

2. Copy your trained model:

```
bash
cp /path/to/best_model.keras models/
```

### 3. Ensure database is in repository:

- The file `data/pm25_data.db` should be in your repo
- Make sure it's NOT in `.gitignore`

## Step 2: Get API Keys

### 1. Taiwan EPA API Key:

- Visit: <https://data.moenv.gov.tw/>
- Register account → Get API key

### 2. Google Gemini API Key:

- Visit: <https://makersuite.google.com/app/apikey>
- Create API key

## Step 3: Railway Setup

### 3.1 Create Railway Account

1. Go to <https://railway.app>
2. Sign up with GitHub (required for verification)
3. Verify your account to enable Full Trial

### 3.2 Create New Project

1. Click "New Project"
2. Select "Deploy from GitHub repo"
3. Connect your GitHub account
4. Select your PM2.5 project repository

### 3.3 Configure Environment Variables

In Railway dashboard → **Variables** tab, add:

env

```
EPA_API_KEY=your_epa_api_key_here  
GEMINI_API_KEY=your_gemini_api_key_here  
SECRET_KEY=your_random_secret_key  
DATABASE_PATH=/data/pm25_data.db  
MODEL_PATH=./models/best_model.keras  
SITE_NAME=板橋  
TZ=Asia/Taipei  
FLASK_ENV=production  
RAILWAY_RUN_UID=0
```

**Important:** `RAILWAY_RUN_UID=0` is required for volume write permissions.

### 3.4 Create Volume

Railway volumes provide persistent storage that survives container restarts.

1. In Railway dashboard, click "+ New" button
2. Select "**Volume**"
3. Configure:
  - **Mount Path:** `/data`
  - **Name:** `pm25-data` (or any name you prefer)
4. Attach the volume to your service

**Note:** If you can't use Cmd+K shortcut, use the "+ New" button in the Railway UI instead.

### Step 4: Deploy

1. Railway will automatically detect `Dockerfile`
2. Click "**Deploy**" (or push to GitHub to auto-deploy)
3. Wait for build (~3-5 minutes)
4. Monitor build logs for any errors

### Step 5: Generate Public Domain

1. Go to **Settings** → **Networking**
2. Click "**Generate Domain**"
3. Railway will provide a public URL like:

<https://pm25-prediction-project-production.up.railway.app>

## Step 6: Verify Deployment

1. Visit your Railway URL

2. Check **Deploy Logs** tab for:

-  Initializing database from backup...
-  Database initialized!
-  SCHEDULER STARTED
-  Fetching data from EPA API...
-  Model loaded successfully
-  Starting Flask app on port 8080...

3. Test the dashboard:

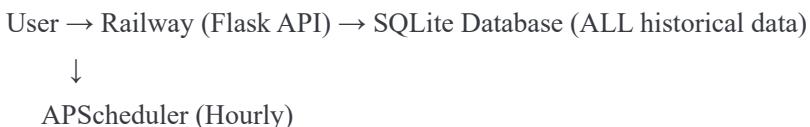
- Current PM2.5 displays
- Historical charts load
- 24-hour predictions show
- Chatbot responds (try: "What was PM2.5 in 2020?")

## How It Works

### Hourly Cycle

- 00:00 → Crawler fetches new data from EPA API  
→ Clean & forward-fill missing values  
→ Insert into SQLite (preserves ALL historical data)  
→ Model: Read last 720 hours → Predict 24 hours  
→ Store predictions in database  
→ Frontend auto-refreshes via JavaScript

### Data Flow



↓

EPA API Crawler → Model Inference (last 720h)

## Model Details

### Architecture Evolution

#### Initial Model: Single-Layer Transformer

- Architecture: 1 Transformer Block
- Embedding Dim: 64
- Feed-Forward Dim: 256
- Total Parameters: 177,560 (~693 KB)
- Validation Loss: ~0.0045

#### Final Model: Dual-Layer Transformer

- Architecture: 2 Stacked Transformer Blocks
- Embedding Dim: 128
- Feed-Forward Dim: 512
- Total Parameters: 1,358,872 (~5.18 MB)
- Validation Loss: 0.00393
- **Performance Improvement:** 15% reduction in validation loss

### Model Components

1. **Projection Layer:** Maps 1D input to 128-dimensional feature space
2. **Positional Encoding:** Sinusoidal encoding for temporal information
3. **Transformer Encoder Block 1:**
  - Multi-Head Attention (8 heads, `key_dim=16`)
  - Feed-Forward Network ( $128 \rightarrow 512 \rightarrow 128$ )
  - Residual Connections & Layer Normalization
4. **Transformer Encoder Block 2:** Same structure as Block 1
5. **Global Average Pooling:** Aggregate temporal features
6. **Dense Output Layer:**  $128 \rightarrow 24$  (predict 24 hours)

## Training Results

### Performance Metrics:

- Overall R<sup>2</sup>: 0.4166
- Overall MAE: 4.60  $\mu\text{g}/\text{m}^3$
- Overall RMSE: 6.33  $\mu\text{g}/\text{m}^3$
- Hour +1 prediction: R<sup>2</sup>=0.80, MAE=2.73  $\mu\text{g}/\text{m}^3$
- Hour +24 prediction: R<sup>2</sup>=0.25, MAE=5.25  $\mu\text{g}/\text{m}^3$

### Training Details:

- Training Time: 258.81 minutes (31 epochs)
- Average: 8.0 min/epoch
- Hardware: Google Colab T4 GPU
- Optimizer: Adam (lr=1e-3)
- Loss Function: MSE
- Early Stopping: Patience=5 epochs

### Input/Output Specification

- **Input:** (batch\_size, 720, 1) - 720 hours of PM2.5 data
- **Output:** (batch\_size, 24) - 24 hours of predictions
- **Lookback Window:** 30 days (720 hours)
- **Forecast Horizon:** 24 hours

### Notes

- System requires at least 720 hours (30 days) of data for predictions
- All historical data preserved (2018-2025, ~61,000+ hours)
- Chatbot can query entire historical dataset via RAG
- Hourly auto-update via EPA API crawler

Good luck! 

