

PM2.5 Air Quality Prediction Dashboard

Real-time PM2.5 monitoring and 24-hour prediction system using Transformer neural network, deployed on Railway with auto-updating hourly crawler.

Features

-  **Real-time PM2.5 Monitoring** - Live data from Taiwan EPA API
 -  **24-Hour Predictions** - Transformer model forecasting
 -  **Interactive Visualizations** - Chart.js powered charts
 -  **AI Chatbot** - Gemini-powered RAG for historical queries
 -  **Auto-Update** - Hourly crawler + model inference
 -  **Modern UI** - Glassmorphism design, light mode
-

Architecture

User → Railway (Flask API) → SQLite Database (ALL historical data)



APScheduler (Hourly)



EPA API Crawler → Model Inference (last 720h)

- **Frontend:** HTML/CSS/JavaScript + Chart.js
 - **Backend:** Flask + APScheduler
 - **Database:** SQLite with Railway Volume (ALL data 2018-2025, ~61,000+ hours)
 - **Model:** Transformer (uses last 720h → predicts 24h)
 - **RAG:** Gemini API (queries full 2018-2025 history)
 - **Deployment:** Railway with Docker
-

Project Structure

pm25_dashboard/

```

└── backend/
    ├── app.py          # Flask main app
    ├── config.py       # Configuration
    ├── database.py     # SQLite utilities
    ├── init_db.py      # Initialize DB from CSV
    ├── crawler.py      # EPA API crawler
    ├── prediction_service.py # Model inference
    ├── rag_service.py  # Gemini chatbot
    └── scheduler.py    # APScheduler

└── frontend/
    ├── index.html
    └── static/
        ├── css/style.css
        └── js/
            ├── dashboard.js
            └── chat.js

└── models/
    └── best_model.keras # YOUR TRAINED MODEL

└── Dockerfile
└── railway.toml
└── requirements.txt
└── .env.example
└── README.md

```

DEPLOYMENT INSTRUCTIONS

Step 1: Prepare Your Data & Model

1. Prepare your full historical CSV:

- Merge all PM2.5 data files (2018-2025) into one CSV
- Format: `(createdAt, pm25)` columns
- Example: 61,000+ rows = ~7 years of hourly data

2. Initialize database with ALL data:

```

bash

# On your local machine
python -m backend.init_db --csv /path/to/all_pm25_data.csv

```

Expected output:

 Data Range:

Start: 2018-01-01 00:00

End: 2025-11-23 14:00

Total: 61,320 hours (2,555 days)

Database initialized successfully!

Total measurements: 61,320

Database size: ~3 MB (estimated)

This creates `data/pm25_data.db` with ALL historical data.

3. Copy your trained model:

```
bash
```

```
cp /path/to/best_model.keras models/
```

Step 2: Get API Keys

1. Taiwan EPA API Key:

- Visit: <https://data.moenv.gov.tw/>
- Register account → Get API key

2. Google Gemini API Key:

- Visit: <https://makersuite.google.com/app/apikey>
- Create API key

Step 3: Railway Setup

3.1 Create Railway Account

1. Go to <https://railway.app>
2. Sign up with GitHub (required for verification)
3. Verify your account to enable **Full Trial**

3.2 Create New Project

1. Click "New Project"
2. Select "Deploy from GitHub repo"

3. Connect your GitHub account
4. Push this project to your GitHub repository
5. Select the repository in Railway

3.3 Configure Environment Variables

In Railway dashboard → **Variables** tab, add:

```
env  
  
EPA_API_KEY=your_epa_api_key_here  
GEMINI_API_KEY=your_gemini_api_key_here  
SECRET_KEY=your_random_secret_key  
DATABASE_PATH=/data/pm25_data.db  
MODEL_PATH=./models/best_model.keras  
SITE_NAME=板橋  
TZ=Asia/Taipei  
FLASK_ENV=production
```

3.4 Create Volume

1. In Railway dashboard, press **K** (Cmd+K)
2. Type "**Volume**" → Select "**New Volume**"
3. Set:
 - **Mount Path:** `/data`
 - **Name:** `pm25-data` (any name)
4. Attach volume to your service

3.5 Set RAILWAY_RUN_UID

Add this variable to fix volume permissions:

```
RAILWAY_RUN_UID=0
```

3.6 Upload Database

Since Railway volumes can't be accessed directly, you have two options:

Option A: Include in Docker Image (Recommended for initial setup)

1. Keep `data/pm25_data.db` in your repository (remove from `.gitignore`)

2. Modify Dockerfile to copy database:

```
dockerfile
```

```
COPY data/pm25_data.db /data/pm25_data.db
```

Option B: Use Railway File Browser Template

1. Deploy the "File Browser" template from Railway
2. Upload `pm25_data.db` to `/data/` directory
3. Delete File Browser service after upload

Step 4: Deploy

1. Railway will automatically detect `Dockerfile`
2. Click "**Deploy**"
3. Wait for build (3-5 minutes)
4. Once deployed, Railway provides a public URL like:

```
https://your-app.railway.app
```

Step 5: Verify Deployment

1. Visit your Railway URL
2. Check console logs in Railway dashboard:

```
⌚ SCHEDULER STARTED
📡 Fetching data from EPA API...
✓ Stored X measurements
🧠 Running model inference...
```

3. Test the dashboard:
 - Current PM2.5 displays
 - Charts load
 - Chatbot responds

Troubleshooting

Database Not Persisting

- Ensure `RAILWAY_RUN_UID=0` is set
- Verify volume mount path is `/data`
- Check logs for permission errors

Model Not Loading

- Ensure `best_model.keras` is in `models/` directory
- Check file size < 100MB (Railway limit)
- Verify `MODEL_PATH` environment variable

Crawler Fails

- Verify `EPA_API_KEY` is correct
- Check EPA API status: <https://data.moenv.gov.tw/>
- Review logs for API errors

Out of Credits

- Railway free trial: \$5 for 30 days
 - After trial: Upgrade to Hobby plan (\$5/month)
 - Monitor usage in Railway dashboard
-

How It Works

Hourly Cycle

```
00:00 → Crawler fetches new data from EPA API
      → Clean & forward-fill missing values
      → Insert into SQLite (preserves ALL historical data)
      → Model: Read last 720 hours → Predict 24 hours
      → Store predictions in database
      → Frontend auto-refreshes via JavaScript
```

Data Flow

1. **Initialization:** ALL historical data (2018-2025) → SQLite
 2. **Hourly Update:** EPA API → New hour → Append to database
 3. **Model Inference:** Last 720 hours → Transformer → 24h forecast
 4. **RAG Queries:** Full database (2018-current) for chatbot
-

Local Development

Setup

```
bash

# Clone repository
git clone <your-repo-url>
cd pm25_dashboard

# Create virtual environment
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate

# Install dependencies
pip install -r requirements.txt

# Create .env file
cp .env.example .env
# Edit .env with your API keys

# Initialize database
python -m backend.init_db --csv data/your_pm25.csv --site 板橋

# Run Flask app
python -m backend.app
```

Access

- Dashboard: <http://localhost:5000>
 - API: <http://localhost:5000/api/current>
-

API Endpoints

Endpoint	Method	Description
/api/current	GET	Current PM2.5 + next hour prediction
/api/predictions	GET	24-hour forecast
/api/history?hours=168	GET	Historical data (default 7 days)
/api/chat	POST	Chatbot query
/api/status	GET	AI health advice
/api/stats	GET	Database statistics

Model Details

- **Architecture:** Dual-layer Transformer
- **Input:** 720 hours (30 days) of PM2.5 data
- **Output:** 24-hour predictions
- **Parameters:** 1.36M (~5.18 MB)
- **Training:** Google Colab T4 GPU
- **Performance:** $R^2=0.42$, MAE=4.60 $\mu\text{g}/\text{m}^3$

Support

- Railway Docs: <https://docs.railway.app>
- EPA API Docs: <https://data.moenv.gov.tw/>
- Gemini API: <https://ai.google.dev/>

License

MIT License - Feel free to use for your projects!

Deployment Checklist

- CSV data merged (2018-2025, all years)
- Database initialized locally (`init_db.py`)
- Model file ready (`best_model.keras`)

- EPA API key obtained
 - Gemini API key obtained
 - GitHub repository created
 - Railway account verified
 - Environment variables configured
 - Volume created and mounted to `/data`
 - `RAILWAY_RUN_UID=0` set
 - Database uploaded to Railway
 - Deployment successful
 - Dashboard accessible via Railway URL
 - Hourly scheduler running
 - RAG queries working (test with "What was PM2.5 in 2020?")
-

Deployment Time: ~20-30 minutes

Good luck! 